

## Clone repository and compiling code

```
# Clone repo
git clone https://github.com/mar-file-system/GUFI.git
# make build directory and change into that directory
mkdir build
cd build
# compiling code
cmake .. -DCMAKE_INSTALL_PREFIX=~/.local/gufi
make -j
sudo make install
# Export to PATH (you can also add it to .bashrc)
export PATH="$PATH:~/.local/gufi/bin/"
```

## Build GUFI Index

```
gufi_dir2index <input dir> <output dir> -n <thread count>
```

## Some Tables in GUFI Index

- `entries`: metadata of the files in the current directory
- `summary`: summary of current directory
- `treесummary`: summary of current and all the subdirectories

## How does aggregation work?

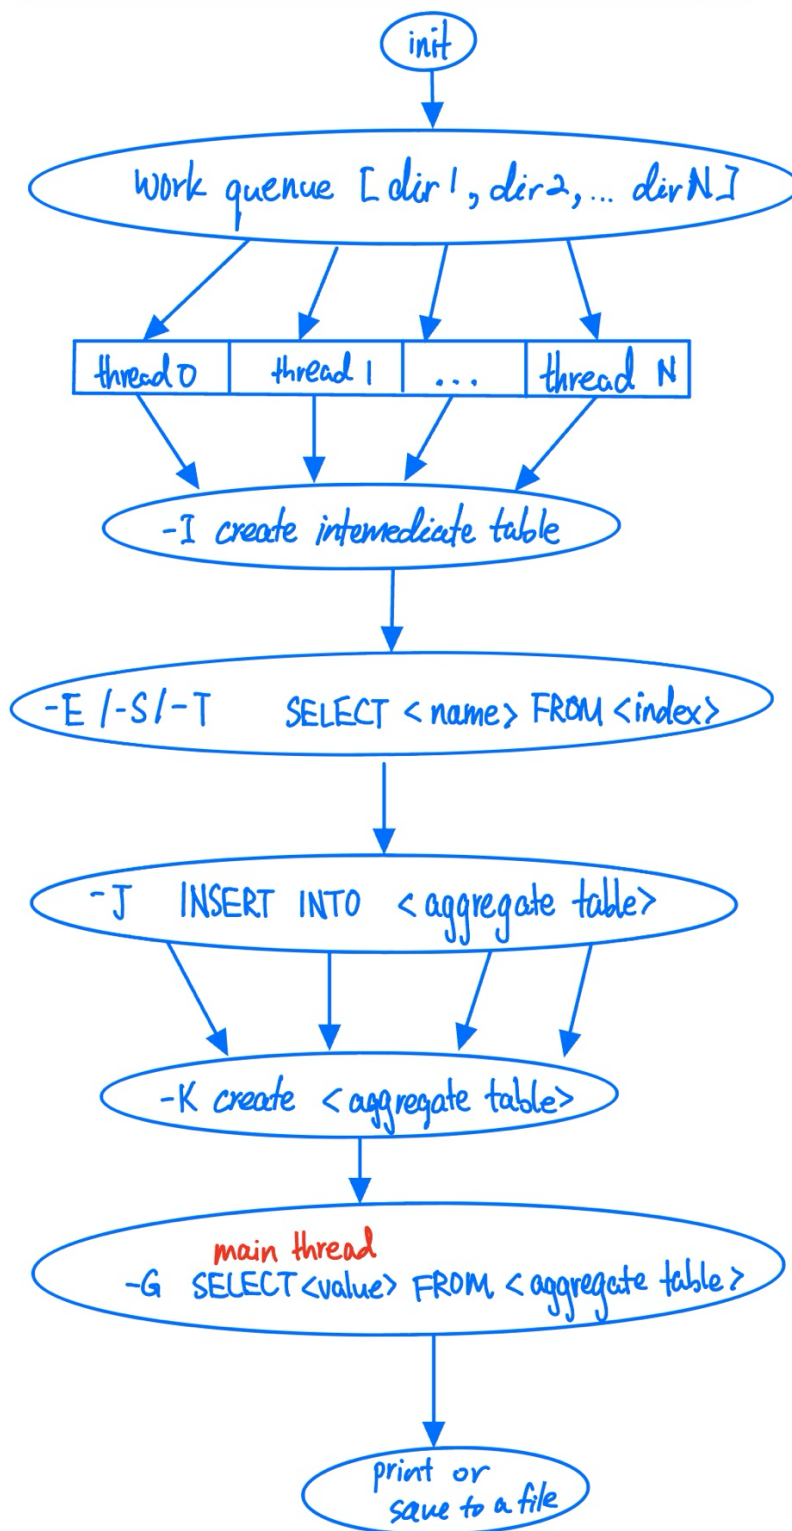
### How does work assignment in (generic) thread pools work?

- `-I`: Create the intermediate table
- `-E / -S / -T`: Each thread executes SQL (e.g., `SELECT from entries/summary`) and writes to `intermediate`
- `-K`: Create the final aggregate table
- `-J`: The main thread aggregates rows from `intermediate` into `aggregate`
- `-G`: The final `SELECT` from `aggregate` to produce user output

### How does this affect `-s/-E` result placement?

- Each thread runs its own of the `-S / -E` query on the directories it processes.
- Results are written into per-thread intermediate table in memory or file, one per thread.
- The data placement depends on which thread processes which directory.
- If using `-o` or `-O`, each thread writes its output to separate files:
  - e.g., `outfile.0`, `outfile.1`, etc.

## Workflow



## gufi\_query Examples

### Running command at Zoom Level D (each directory)

### Zoom Level D

```
graph TD; A[A] --> B[B]; A --> C[C]; B --> D[D]; B --> E[E]; B --> F[F]; E --> I[I]; I --> K[K]; K --> L[L]; K --> M[M]; K --> N[N]; K --> O[O]; C --> G[G]; C --> H[H]; H --> J[J];
```

- View each directory as a collection of data points
  - Size → min/max/mean/stddev/etc. of all files at directory
- More resolution than Zoom Level S
  - Too much detail
    - Currently at ~25-50 Million Directories
    - Target is 1 Billion Directories
- Use `longitudinal_snapshot.py` directory to collect this data

Los Alamos NATIONAL LABORATORY

10/21/2024 29

- -E: for \*entries table, such as pentries, entries
- -S: for \*summary table, such as summary, vrsummary
- -T: for treesummary table

**# Get size from entries table (32 threads) -- size of each file in the index**

```
gufi_query <index> -E "SELECT size FROM entries" -n 32
```

**# Get size from summary table (32 threads) -- size of each directory in the index**

```
gufi_query <index> -S "SELECT size FROM summary" -n 32
```

**# To access treesummary, first needs to create treesummary table**

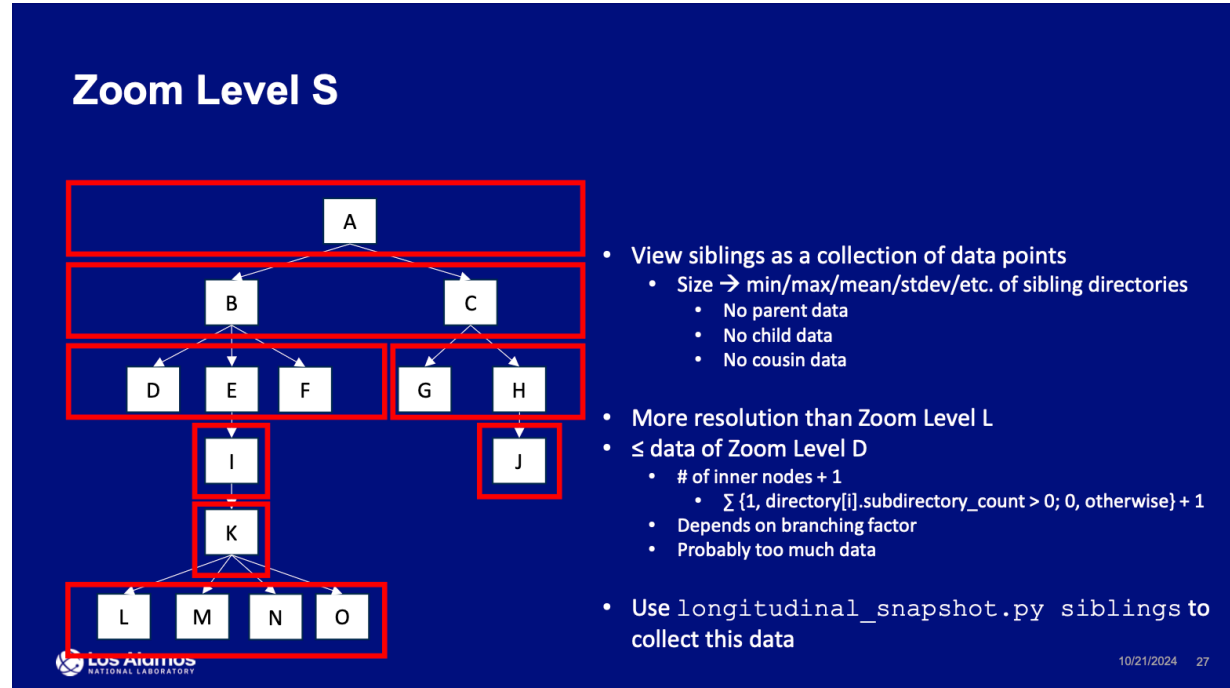
```
gufi_treesummary_all <index> -n 32
```

**# Get size from treesummary table (32 threads) -- size of each directory + all subdirectory in the index**

**# in the example image, if we are getting size at directory B, it will also includes all the size below (D, E, F, I, K, L, M, N, O)**

```
gufi_query <index> -S "SELECT size FROM treesummary" -n 32
```

## Running command at Zoom Level S (siblings)



```
# Get size from pentries table (32 threads)
```

```
# -I creates an intermediate table to store per thread level,  
ppinode (to get the file from same direcorry) and sum of size for  
-E
```

```
# -E gets file's level, ppinode and sum of size from the entries  
table insert into intermediate table
```

```
# -K creates an aggregate table to combine per thread level,  
ppinode and size into one dataset
```

```
# -J gets level and sum of size (of matching levels and ppinode)  
from intermediate table per thread combine into one data store  
into aggregate table
```

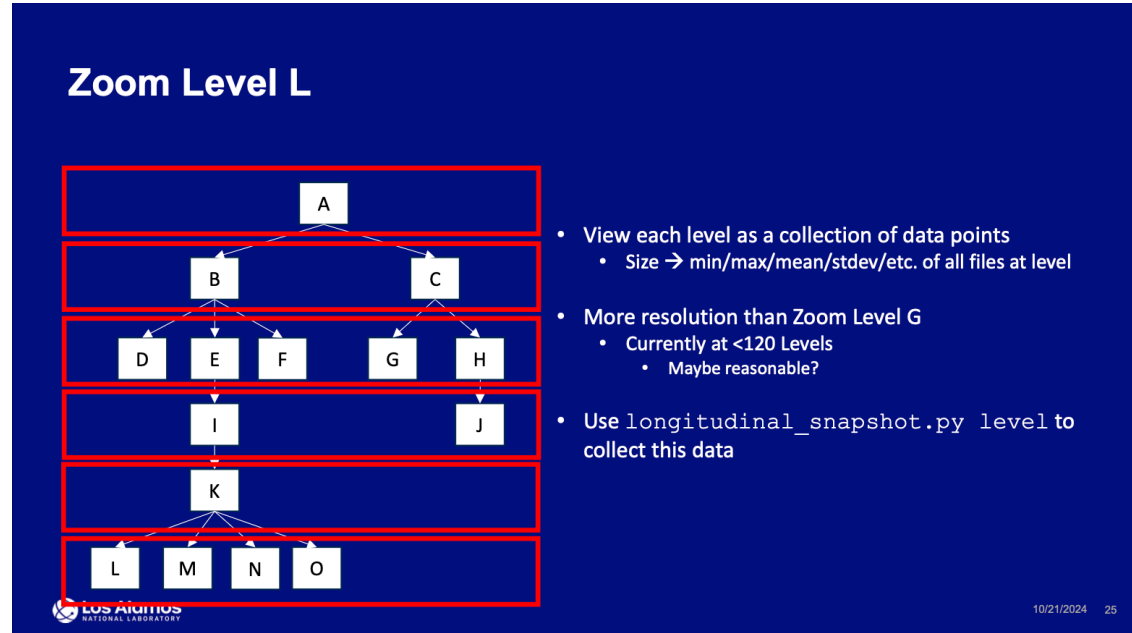
```
# at this point, there are level, ppinode and sum of size from  
each thread potentially with repeated levels and ppinode  
originating from different thread
```

```
# -G groups them to remove duplicate ppinode and sum to get down  
to one size and print
```

```
# -d is the 1-character separate for each column
```

```
gufi_query <index> \  
-I "CREATE TABLE intermediate(level INT64, ppinode TEXT,  
total_size INT64);" \  
-E "INSERT INTO intermediate SELECT level() AS level, ppinode,  
SUM(size) FROM pentries WHERE type = 'f' GROUP BY level(),  
ppinode;" \  
-K "CREATE TABLE aggregate(level INT64, ppinode TEXT, total_size  
INT64);" \  
-J "INSERT INTO aggregate SELECT level, ppinode, SUM(total_size)  
FROM intermediate GROUP BY level, ppinode;" \  
-G "SELECT level, ppinode, SUM(total_size) FROM aggregate GROUP  
BY level, ppinode;" \  
-d '|' -n 32
```


## Running command at Zoom Level L (level)



```
# -I creates an intermediate table to store per thread level and
sum of size for -E
# -E gets file's level and sum of size from the entries table
insert into intermediate table
# -K creates an aggregate table to combine per thread level and
size into one dataset
# -J gets level and sum of size (of matching levels) from
intermediate table per thread combine into one data store into
aggregate table
# at this point, there are level and sum of size from each
thread potentially with repeated levels originating from
different thread
# -G groups them to remove duplicate levels and sum to get down
to one size and print
# -d is the 1-character separate for each column
gufi_query \
-I "CREATE TABLE intermediate(level INT64, total_size INT64);" \
-E "INSERT INTO intermediate SELECT level() AS level, SUM(size)
FROM entries WHERE type = 'f';" \
-K "CREATE TABLE aggregate(level INT64, total_size INT64);" \
-J "INSERT INTO aggregate SELECT level, SUM(total_size) FROM
intermediate GROUP BY level;" \
-G "SELECT level, SUM(total_size) FROM aggregate GROUP BY
level;" \
-d ' ' -n 32 <index>
```

## Running command at Zoom Level G (global)

### Zoom Level G



- View entire tree as one collection of data points
  - Size → min/max/mean/stddev/etc. of all files
- Low resolution
- Use `longitudinal_snapshot.py` graph to collect this data
- Treesummary data can be collected, but is redundant

```
# -I creates an intermediate table to store per thread data for
-E
# -E gets file's size from the entries table insert into
intermediate table
# -K creates an aggregate table to combine per thread size into
one dataset
# -J gets sum of size from intermediate table per thread combine
into one data store into aggregate table
# at this point, there is one sum of size in the aggregate table
# -G get size and print
# -d is the 1-character separate for each column
gufi_query \
-I "CREATE TABLE intermediate(size INT64);" \
-E "INSERT INTO intermediate SELECT size FROM entries WHERE type
= 'f';" \
-K "CREATE TABLE aggregate(total INT64);" \
-J "INSERT INTO aggregate SELECT SUM(size) FROM intermediate;" \
-G "SELECT total FROM aggregate;" \
-d '|' -n 32 <index>
```